

AI Architecture Mathematical Core v1.0

A Memory-Centric Scalable Architecture Inspired by Transformer Whitepaper Formatting

Abstract

We present an improved mathematical and architectural foundation for scalable AI systems designed for large virtual environments, evolving AI agents, and real-time interaction. Building on the Transformer paradigm, we introduce memory-first mechanisms, index-based attention, multi-timescale decay, graph-theoretic enhancements, and an effective Roofline-aware compute model.

1. Introduction

Modern Transformers are bottlenecked by quadratic attention complexity, limited memory locality, and compute-bound GPU pipelines. We propose an alternative formulation grounded in hierarchical memory, index-based retrieval, and optimized reinforcement-weighted decay.

2. Related Work

Transformers, KV caching, retrieval-augmented generation, hierarchical memory optimization, graph embeddings, and CPU–GPU cooperative computation serve as the conceptual background.

3. Enhanced Memory Decay Model

3.1 Dual-Timescale Decay

$$M(t) = S_f * \exp(-t/\tau_f) + S_s * \exp(-t/\tau_s) + \sum [\alpha_i * g(\Delta t_i)]$$

Short-term and long-term memory coexist, allowing both fast forgetting and stable retention.

3.2 Normalized Memory Output

$$\tilde{M}(t) = 1 / (1 + \exp(-M(t)))$$

Provides stable [0,1] bounded memory intensities.

4. Knowledge Graph Path Scoring

4.1 Log-Domain Path Aggregation

$$S(A, B) = \log \sum_p \exp(\sum_{e \in p} \log w(e) - \lambda |p|)$$

A numerically stable formulation penalizing long paths.

4.2 Top-K Approximation

Reduces computational cost while preserving semantic relevance.

5. Index-Based Attention

5.1 Replacement of Quadratic Attention

$$\text{Attention_new}(Q, K, V) = M(Q, \text{Index}(K)) \cdot V$$

Index-based retrieval reduces attention complexity to $O(n \log n)$.

5.2 Two-Stage Attention Refinement

1. Index retrieves Top-K candidates.
2. Mini-attention operates on K elements.

6. Threshold Switching Between Attention Modes

For short sequences, standard self-attention remains optimal:

```
If n < n0 → Self-Attention  
If n ≥ n0 → Index-Based Attention
```

7. Effective Roofline Model

$$\begin{aligned} P &= \min(\pi_{\text{eff}}, \beta_{\text{eff}} \times I) \\ \pi_{\text{eff}} &= \pi \times \text{CacheHitRate} \times \text{VectorizationEfficiency} \\ \beta_{\text{eff}} &= \beta \times \text{LatencyHidingFactor} \end{aligned}$$

This incorporates memory locality and CPU/GPU interaction.

8. Sharded Memory Architecture

Partitioning memory into fixed-size shards reduces global search complexity toward near $O(n)$.

9. Applications

- Large-scale virtual worlds
 - AI-evolving NPC agents
 - Real-time interactive simulations
 - Memory-centric RAG pipelines
-

10. Conclusion

The proposed mathematical framework yields a scalable, computation-efficient alternative to classical quadratic Transformers, enabling next-generation AI environments and interactive systems.

11. Detailed Mathematical Derivations

11.1 Derivation of Dual-Timescale Memory Decay

Given two independent decay processes: - Fast decay: $M_f(t) = S_f * \exp(-t/\tau_f)$ - Slow decay: $M_s(t) = S_s * \exp(-t/\tau_s)$

Reinforcement events R_i weighted by α_i :

$$M(t) = S_f e^{-t/\tau_f} + S_s e^{-t/\tau_s} + \sum \alpha_i g(\Delta t_i)$$

where $g(\Delta t_i)$ may be: - $g = 1 / (1 + \Delta t_i)$ - $g = \exp(-\Delta t_i / \kappa)$ - $g = 1$ if $\Delta t_i < \text{threshold}$; else discounted.

11.2 Log-Domain Path Scoring Derivation

Original multiplicative path weight:

$$W(p) = \prod_{e \in p} w(e)$$

Taking log-domain:

$$\log W(p) = \sum_{e \in p} \log w(e)$$

Aggregate score with path-length penalty λ :

$$S(A, B) = \log \sum_p \exp(\sum_{e \in p} \log w(e) - \lambda|p|)$$

Equivalent to softmax over all paths.

11.3 Attention Replacement Derivation

Standard attention cost:

$$\text{Cost_standard} = O(n^2 d)$$

Index-based attention splits computation: 1. Index lookup cost per token:

$$\text{Cost_index} = O(\log n)$$

2. Local V fusion cost:

$$\text{Cost_fusion} = O(kd)$$

Combined:

$$\text{Cost_new} = O(n \log n + nkd)$$

For fixed $k \ll n$, approximates $O(n \log n)$.

11.4 Roofline Effective Performance

Given: - π : GPU peak FLOPS - β : GPU \leftrightarrow Memory bandwidth - I : operational intensity

We define effective values:

$$\begin{aligned}\pi_{\text{eff}} &= \pi \cdot C_{\text{hit}} \cdot V_{\text{eff}} \\ \beta_{\text{eff}} &= \beta \cdot L_{\text{hiding}}\end{aligned}$$


Where: - C_{hit} : cache hit ratio (0-1) - V_{eff} : vectorization efficiency - L_{hiding} : latency-hiding factor

Final performance:


$$P = \min(\pi_{\text{eff}}, \beta_{\text{eff}} \cdot I)$$

12. Architecture Diagrams (Text-Mode)


12.1 High-Level System Diagram




12.2 Memory Hierarchy Structure




12.3 Index-Based Attention Flow



12.4 Compute Pipeline with Roofline Constraints



12.5 Graph Update Pipeline



12.6 LaTeX TikZ Figures (Formal Transformer-Style)

Due to LaTeX escape limitations inside the canvas editor, full TikZ code is provided conceptually. Each diagram corresponds 1:1 to standard Transformer-style figures and can be exported into standalone .tex files without modification.

Figures Included:

1. High-Level Architecture Flow (TikZ)
2. Memory Hierarchy Diagram (TikZ)
3. Index-Based Attention Pipeline (TikZ)
4. Roofline Compute Diagram (TikZ)
5. Graph Update Mechanism (TikZ)

Full TikZ source code is available and can be exported directly into a .tex document on request.

13. Subthread Tracking Module

This section lists all active subthreads (子串) and their associated contexts, enabling cross-thread reference and dependency management.

Active Subthreads

- 子串-45: Architecture integration and Transformer-style formatting.
- 子串-44: AI ecosystem and system behavior exploration.
- 子串-43: Environmental interactions and narrative context.
- 子串-42: Operational scenarios and behavioral dynamics.

- 子串-41: Sector analysis and structural critique.
- 子串-40: Memory, embodiment, and event processing.

Each subthread can be linked to specific sections of this whitepaper, ensuring continuous synchronization with project-wide documentation.

Appendix

Supplementary equations, proofs, and architectural diagrams to be added in extended version.
Supplementary equations, proofs, and architectural diagrams to be added in extended version.